

DIFUSCO: Graph-based Diffusion Solvers for Combinatorial Optimization

2025. 04. 11

AI&OPT
김정현


Contents

1. Introduction
2. Related Work
3. Diffusion Model
4. Proposed Method (DIFUSCO)
5. Experiment
6. Result
7. Conclusion


01 Introduction

❖ Combinatorial Optimization (CO) Problem

- ✓ A mathematical problem that aims to find the optimal solution in a discrete solution space.
- ✓ Often belongs to the class of NP-Complete (NPC) problems, making them difficult to solve in polynomial time.
- ✓ Traditionally solved using integer programming and heuristic-based methods.
- ✓ Examples :
 - TSP (Traveling Salesman Problem)
: Finding the shortest possible route that visits each city exactly once and returns to the starting point.
 - MIS (Maximum Independent Set) :
: Finding the largest subset of nodes in a graph such that no two nodes are connected.



<TSP>



<MIS>

02 Related Work

❖ Autoregressive Constructive Solvers

- ✓ Generate solutions sequentially, one element at a time.
- ✓ High computational complexity
- ✓ Difficult to scale to large problem sizes

❖ Non-Autoregressive Constructive Solvers

- ✓ Generate the entire solution in one shot.
- ✓ Struggles to model multi-modal distributions
- ✓ Hard to find good solutions when multiple optima exist (e.g., in the same graph)

❖ Improvement Heuristic Solvers

- ✓ Iteratively refine an initial solution.
- ✓ Examples: 2-opt, node swap
- ✓ Sparse rewards and low sampling efficiency → leads to slow training and inference


03 Diffusion Model

❖ Problem Definition

- ✓ For a given problem instance s , the set of feasible solutions is represented as a binary vector space $x \in \{0,1\}^n$.
 - TSP: Whether each edge is selected
 - MIS: Whether each node is selected
- ✓ The objective is to minimize the following cost: $c_s(x) = cost(x, s) + valid(x, s)$
 - $cost_{TSP}(x, s) = \sum_i x_i \cdot d_i^{(s)}$ (sum of selected edge distances)
 - $cost_{MIS}(x, s) = \sum_i (1 - x_i)$ (maximize number of selected independent nodes)

❖ Diffusion Model

- ✓ Originally used for image and text generation, diffusion models are now being applied to CO problems.
- ✓ Learns to generate optimal solutions via supervised training.
- ✓ Each training sample provides optimal solution x_s^* → Train the model to generate x_s^* as accurately as possible.
- ✓ Forward process: add noise to clean solution $x_0 \rightarrow x_T$
- ✓ Reverse process: Denoise x_T step-by-step to recover x_0
- ✓ Types of Noise:
 - Continuous (Gaussian) : Used for real-valued data
 - Discrete (Bernoulli) : Used for binary solution spaces



04 DIFUSCO_Dataset preparation

❖ Data Generation

- ✓ Randomly generate N cities as 2D coordinates (x, y) in the range $[0, 1]$.
- ✓ Use LKH or Concorde solver to compute the optimal tour.
- ✓ Convert the tour into a binary adjacency matrix.
- ✓ Save each instance to .txt file:
 - $[x_1\ y_1\ x_2\ y_2\ \dots\ x_N\ y_N\ \text{output}\ t_1\ t_2\ \dots\ t_N\ t_1]$
 - t_1 is repeated to form a closed loop.

❖ Dataset Construction

- ✓ Load the .txt file line by line
- ✓ For each line:
 - Extract coordinates \rightarrow points
 - Extract tour \rightarrow build adj_matrix

❖ Batching & Model Input

- ✓ DataLoader creates mini-batches (e.g., batch size = 128).
- ✓ Each batch includes: points, adj_matrix, t

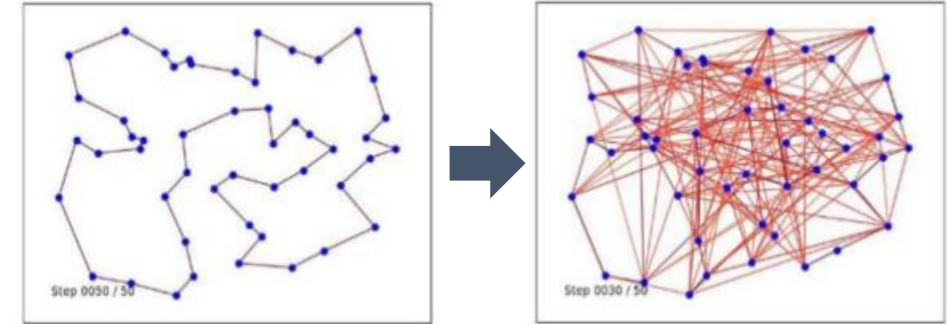
04 DIFUSCO_Forward process

❖ Objective

- ✓ To create a noisy version x_t from the original solution x_0 .
- ✓ So, the model can learn to denoise it back during training.

❖ Step-by-step process

- ✓ From the ground-truth TSP tour, construct the binary adjacency matrix x_0 .
- ✓ Convert each edge(0 or 1) into a one-hot vector.
- ✓ Randomly sample a timestep $t \in [1, T]$.
- ✓ Using the cumulative transition matrix Q_t , add noise to x_0 to generate a noised sample x_t .
- ✓ For training stability:
 - The binary sample $x_t \in \{0,1\}$ is scaled to the range $[-1, 1]$.
 - Small random noise (5%) is added to further regularize the input.



| Tour | | | | | Adj_matrix (x0) | | | | | xt | | | | | [-1, 1] scale | | | | | + small random noise → final xt | | | | | |
|------|---|---|---|---|-----------------|---|---|---|---|----|---|---|---|---|---------------|----|----|----|----|---------------------------------|-------|-------|-------|-------|-------|
| 0 | 1 | 4 | 3 | 2 | 0 | 1 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 1 | -1 | 1 | -1 | -1 | 1 | -1.03 | 1.01 | -1.04 | -1.03 | 1.04 |
| | | | | | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | 1 | 1 | 1 | 1 | -1 | -1 | 1 | 1 | 1.03 | 1.00 | -1.01 | -1.02 | -1.00 |
| | | | | | 1 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | 0 | -1 | 1 | 1 | -1 | -1 | -1 | 1.02 | 1.01 | -1.03 | -1.02 | -1.01 |
| | | | | | 0 | 0 | 1 | 0 | 0 | 1 | 1 | 1 | 0 | 0 | -1 | 1 | 1 | 1 | -1 | -1 | 1.01 | 1.01 | 1.03 | -1.04 | -1.01 |
| | | | | | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 0 | 0 | 1 | 1 | -1 | -1 | -1 | 1 | 1 | 1.00 | -1.04 | -1.03 | -1.01 | 1.02 |

04 DIFUSCO_Reverse process

❖ Objective

- ✓ To recover the original clean adjacency matrix x_0 from a noisy version x_t .

❖ Graph-based Denoising Network (AGNN)

- ✓ Input Components:
 - Node coordinates(x), Noisy adjacency matrix(x_t), current diffusion time step (t)
- ✓ Embedding Steps:
 - Node Embedding: Node, edge, and timestep embeddings.
 - Edge Gating:
 - Decide how much info node i receives from node j .
 - Message Passing:
 - Aggregate messages and update node features.
- ✓ Output: x_0_pred → edge logits
 - Represents predicted probability of edge presence.

❖ Loss

- ✓ Use cross entropy loss with ground-truth x_0 .
- ✓ Predicts whether edge exists (1) or not (0)

| x | | x _t | | | | |
|------|------|----------------|-------|-------|-------|-------|
| 0.94 | 0.41 | -1.03 | 1.01 | -1.04 | -1.03 | 1.04 |
| 0.89 | 0.63 | 1.03 | 1.00 | -1.01 | -1.02 | -1.00 |
| 0.40 | 0.79 | 1.02 | 1.01 | -1.03 | -1.02 | -1.01 |
| 0.67 | 0.02 | 1.01 | 1.01 | 1.03 | -1.04 | -1.01 |
| 0.95 | 0.91 | 1.00 | -1.04 | -1.03 | -1.01 | 1.02 |

x_0_pred

| | | | | | | | | | |
|------|------|------|------|------|-------|-------|-------|-------|-------|
| 0.75 | 0.75 | 0.67 | 0.76 | 0.74 | -0.47 | -0.47 | -0.26 | -0.47 | -0.47 |
| 0.77 | 0.77 | 0.67 | 0.68 | 0.75 | -0.48 | -0.48 | -0.26 | -0.48 | -0.47 |
| 0.67 | 0.75 | 0.67 | 0.66 | 0.66 | -0.25 | -0.47 | -0.27 | -0.26 | -0.37 |
| 0.74 | 0.75 | 0.74 | 0.75 | 0.66 | -0.43 | -0.44 | -0.47 | -0.48 | -0.26 |
| 0.75 | 0.74 | 0.76 | 0.75 | 0.66 | -0.42 | -0.47 | -0.46 | -0.48 | -0.37 |

x_0

| | | | | |
|---|---|---|---|---|
| 0 | 1 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 1 |
| 1 | 0 | 0 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 1 | 0 |

Cross Entropy Loss
= 0.4961

04 DIFUSCO_Inference

❖ Objective

- ✓ To reconstruct the original clean solution x_0 from a fully noised sample x_T .

❖ Inference Process

- ✓ Initialization: Start from a fully noised adjacency matrix x_T .
- ✓ Iterative Denoising Loop (Reverse Diffusion):
 - Model predicts x_{0_pred} from current x_t .
 - Compute the posterior $q(x_{t-1} | x_t, x_{0_pred})$ using diffusion rule.
 - Sample x_{t-1} from the posterior.
 - Repeat until reaching x_0 .
- ✓ Output: The final x_0 is a denoised adjacency matrix representing edge probabilities.
- ✓ May contain multiple disconnected subtours → requires post-processing

❖ Post-processing

- ✓ Merge disconnected subtours into a valid complete tour
- ✓ Apply 2-opt algorithm to refine the path
- ✓ Compute the final tour cost using Euclidean distance.

04 DIFUSCO_Inference

❖ Fast Inference Scheduling

- ✓ To reduce inference time, skip full diffusion steps:
 - Use only a subset of the total steps (e.g., 50 out of 1000)
- ✓ Step selection strategies:
 - Linear: Uniform step intervals
 - Cosine: Large intervals early, denser near the end

❖ Decoding Strategy

- ✓ At the final step, the diffusion model outputs a heatmap
- ✓ Heatmap represents the confidence or probability of each variable (node or edge) being part of the optimal solution
- ✓ TSP:
 - Edge heatmap: Confidence of each edge being part of the tour
 - Decoding methods: Greedy + 2-opt, Sampling, MCTS
- ✓ MIS:
 - Node heatmap: Confidence of each node being in the independent set
 - Select non-conflicting nodes greedily based on ranking

05 Experiment

❖ Datasets

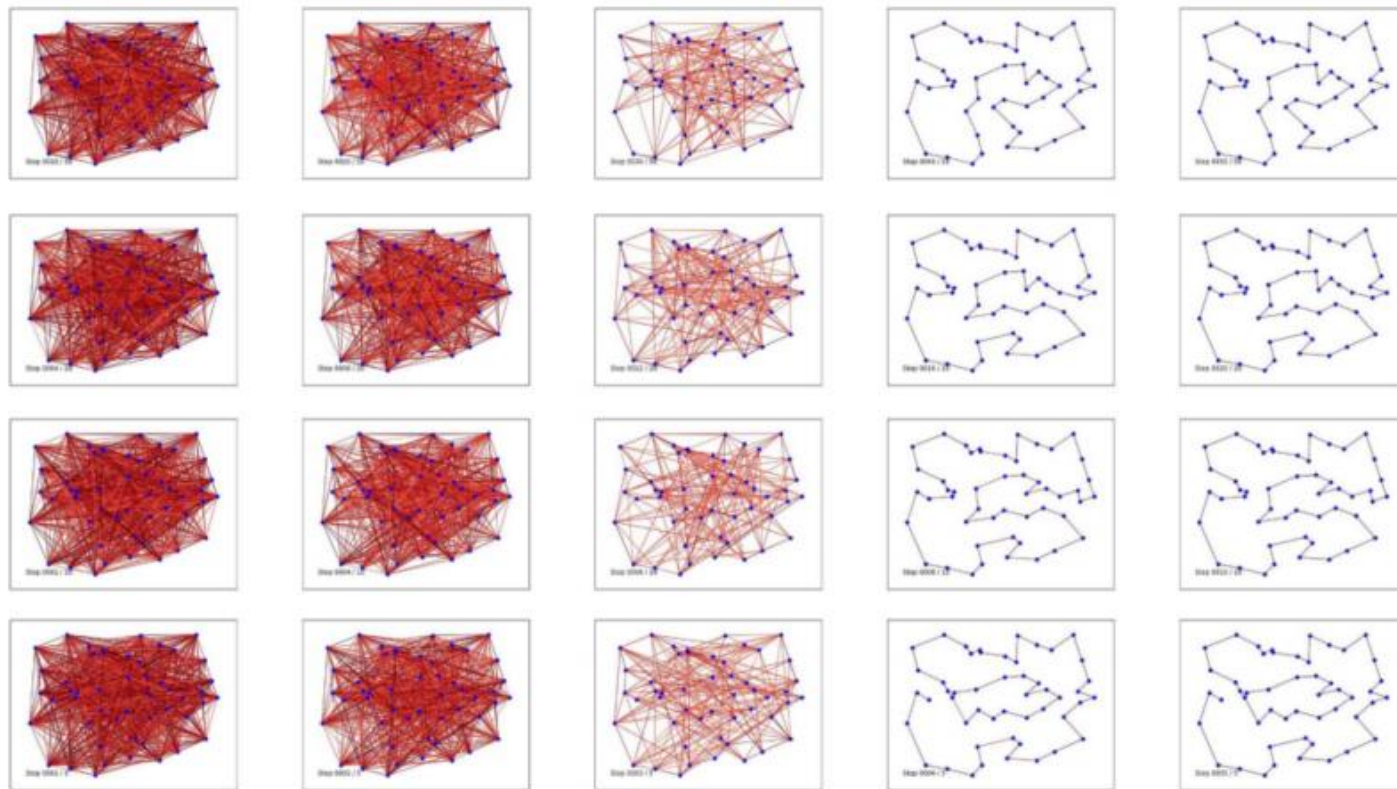
- ✓ TSP50 / TSP100 → Training labels generated using the Concorde exact solver
- ✓ TSP500 / TSP1000 / TSP 10000 → Training labels generated using the LKH-3 heuristic solver
- ✓ Sparse graph: Edge connections are limited to reduce computational complexity

❖ Model Settings

- ✓ Denoising step : 1000
- ✓ Linear noise schedule
- ✓ Decoding Strategy : Greedy + 2-opt

❖ Evaluation Metrics

- ✓ Tour length
- ✓ Performance gap
- ✓ Run time



05 Experiment

❖ **DIFUSCO training setting**

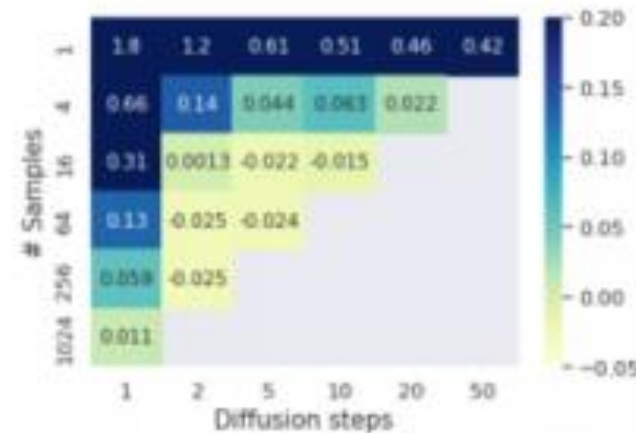
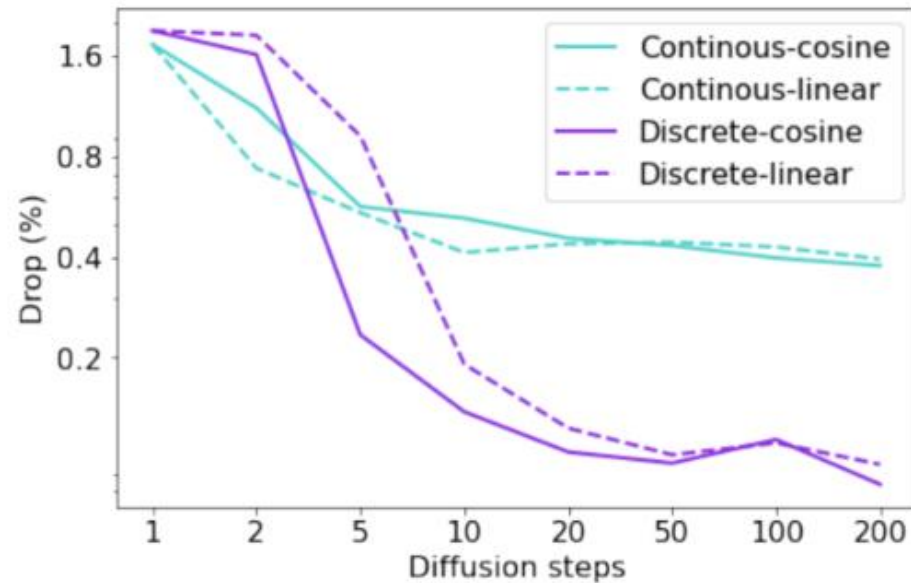
- Curriculum Learning (TSP-500,1000,10000)

| Problem | Dataset | Instances | Epochs | Batch Size |
|---------|--------------|-----------|--------|------------|
| TSP | TSP-50 | 1,502,000 | 50 | 512 |
| | TSP-100 | 1,502,000 | 50 | 256 |
| | TSP-500 | 128,000 | 50 | 64 |
| | TSP-1000 | 64,000 | 50 | 64 |
| | TSP-10000 | 6,400 | 50 | 8 |
| MIS | SATLIB | 49,500 | 50 | 128 |
| | ER-[700,800] | 163,840 | 50 | 32 |

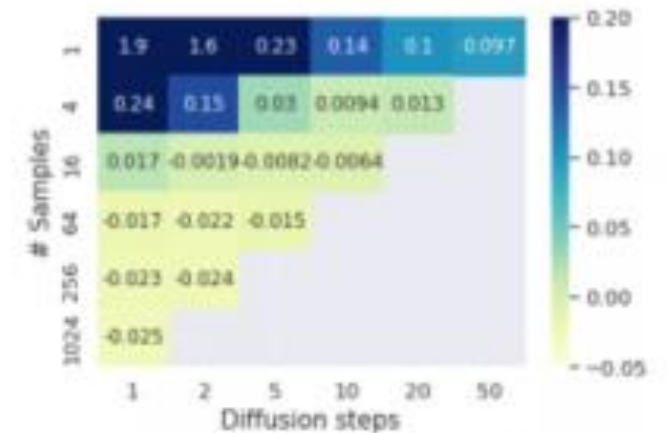
06 Result

❖ Design Analysis

- ✓ Discrete vs Continuous
 - Discrete diffusion with cosine scheduling shows better performance compared to continuous variants
- ✓ More diffusion iterations vs More sampling
 - Results suggest that increased sampling with fewer steps can maintain performance while reducing inference time
 - 50 (diffusion steps) X 1 (samples) and 10 (diffusion steps) X 16 (samples)



(a) Continuous diffusion



(b) Discrete diffusion

06 Result

❖ Comparison to SOTA methods

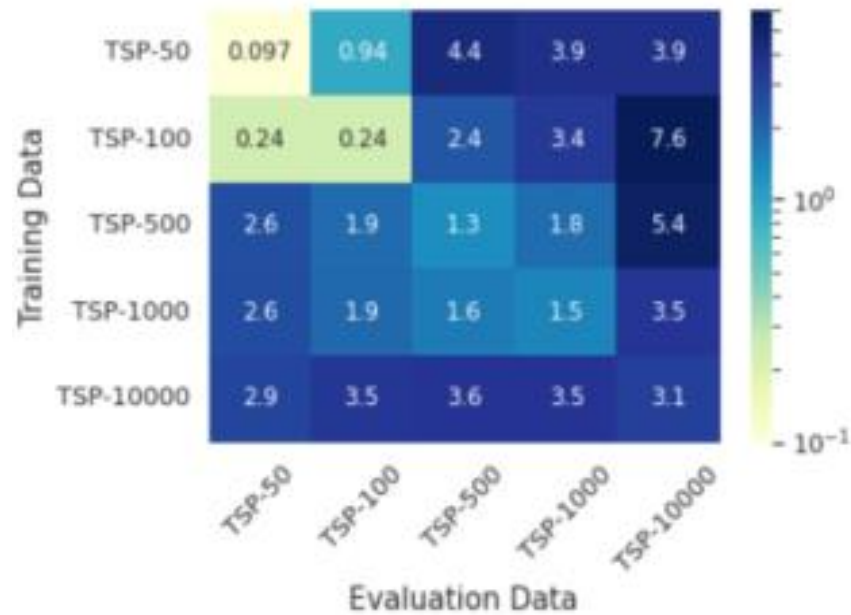
| ALGORITHM | TYPE | TSP-50 | | TSP-100 | |
|-----------------|---------------------|-------------|--------------|-------------|--------------|
| | | LENGTH↓ | GAP(%)↓ | LENGTH↓ | GAP(%)↓ |
| CONCORDE* | EXACT | 5.69 | 0.00 | 7.76 | 0.00 |
| 2-OPT | HEURISTICS | 5.86 | 2.95 | 8.03 | 3.54 |
| AM | GREEDY | 5.80 | 1.76 | 8.12 | 4.53 |
| GCN | GREEDY | 5.87 | 3.10 | 8.41 | 8.38 |
| TRANSFORMER | GREEDY | 5.71 | 0.31 | 7.88 | 1.42 |
| POMO | GREEDY | 5.73 | 0.64 | 7.84 | 1.07 |
| SYM-NCO | GREEDY | - | - | 7.84 | 0.94 |
| DPDP | 1k-IMPROVEMENTS | 5.70 | 0.14 | 7.89 | 1.62 |
| IMAGE DIFFUSION | GREEDY [†] | 5.76 | 1.23 | 7.92 | 2.11 |
| OURS | GREEDY [†] | 5.70 | 0.10 | 7.78 | 0.24 |
| AM | 1k×SAMPLING | 5.73 | 0.52 | 7.94 | 2.26 |
| GCN | 2k×SAMPLING | 5.70 | 0.01 | 7.87 | 1.39 |
| TRANSFORMER | 2k×SAMPLING | 5.69 | 0.00 | 7.76 | 0.39 |
| POMO | 8×AUGMENT | 5.69 | 0.03 | 7.77 | 0.14 |
| SYM-NCO | 100×SAMPLING | - | - | 7.79 | 0.39 |
| MDAM | 50×SAMPLING | 5.70 | 0.03 | 7.79 | 0.38 |
| DPDP | 100k-IMPROVEMENTS | 5.70 | 0.00 | 7.77 | 0.00 |
| OURS | 16×SAMPLING | 5.69 | -0.01 | 7.76 | -0.01 |

| ALGORITHM | TYPE | TSP-500 | | | TSP-1000 | | | TSP-10000 | | |
|---------------------|--------------------------|--------------|--------------|---------------|--------------|--------------|---------------|--------------|--------------|---------------|
| | | LENGTH↓ | GAP↓ | TIME↓ | LENGTH↓ | GAP↓ | TIME↓ | LENGTH↓ | GAP↓ | TIME↓ |
| CONCORDE | EXACT | 16.55* | — | 37.66m | 23.12* | — | 6.65h | N/A | N/A | N/A |
| GUROBI | EXACT | 16.55 | 0.00% | 45.63h | N/A | N/A | N/A | N/A | N/A | N/A |
| LKH-3 (DEFAULT) | HEURISTICS | 16.55 | 0.00% | 46.28m | 23.12 | 0.00% | 2.57h | 71.77* | — | 8.8h |
| LKH-3 (LESS TRAILS) | HEURISTICS | 16.55 | 0.00% | 3.03m | 23.12 | 0.00% | 7.73m | 71.79 | — | 51.27m |
| FARTHEST INSERTION | HEURISTICS | 18.30 | 10.57% | 0s | 25.72 | 11.25% | 0s | 80.59 | 12.29% | 6s |
| AM | RL+G | 20.02 | 20.99% | 1.51m | 31.15 | 34.75% | 3.18m | 141.68 | 97.39% | 5.99m |
| GCN | SL+G | 29.72 | 79.61% | 6.67m | 48.62 | 110.29% | 28.52m | N/A | N/A | N/A |
| POMO+EAS-EMB | RL+AS+G | 19.24 | 16.25% | 12.80h | N/A | N/A | N/A | N/A | N/A | N/A |
| POMO+EAS-TAB | RL+AS+G | 24.54 | 48.22% | 11.61h | 49.56 | 114.36% | 63.45h | N/A | N/A | N/A |
| DIMES | RL+G | 18.93 | 14.38% | 0.97m | 26.58 | 14.97% | 2.08m | 86.44 | 20.44% | 4.65m |
| DIMES | RL+AS+G | 17.81 | 7.61% | 2.10h | 24.91 | 7.74% | 4.49h | 80.45 | 12.09% | 3.07h |
| OURS (DIFUSCO) | SL+G [†] | 18.35 | 10.85% | 3.61m | 26.14 | 13.06% | 11.86m | 98.15 | 36.75% | 28.51m |
| OURS (DIFUSCO) | SL+G [†] +2-OPT | 16.80 | 1.49% | 3.65m | 23.56 | 1.90% | 12.06m | 73.99 | 3.10% | 35.38m |
| EAN | RL+S+2-OPT | 23.75 | 43.57% | 57.76m | 47.73 | 106.46% | 5.39h | N/A | N/A | N/A |
| AM | RL+BS | 19.53 | 18.03% | 21.99m | 29.90 | 29.23% | 1.64h | 129.40 | 80.28% | 1.81h |
| GCN | SL+BS | 30.37 | 83.55% | 38.02m | 51.26 | 121.73% | 51.67m | N/A | N/A | N/A |
| DIMES | RL+S | 18.84 | 13.84% | 1.06m | 26.36 | 14.01% | 2.38m | 85.75 | 19.48% | 4.80m |
| DIMES | RL+AS+S | 17.80 | 7.55% | 2.11h | 24.89 | 7.70% | 4.53h | 80.42 | 12.05% | 3.12h |
| OURS (DIFUSCO) | SL+S | 17.23 | 4.08% | 11.02m | 25.19 | 8.95% | 46.08m | 95.52 | 33.09% | 6.59h |
| OURS (DIFUSCO) | SL+S+2-OPT | 16.65 | 0.57% | 11.46m | 23.45 | 1.43% | 48.09m | 73.89 | 2.95% | 6.72h |
| ATT-GCN | SL+MCTS | 16.97 | 2.54% | 2.20m | 23.86 | 3.22% | 4.10m | 74.93 | 4.39% | 21.49m |
| DIMES | RL+MCTS | 16.87 | 1.93% | 2.92m | 23.73 | 2.64% | 6.87m | 74.63 | 3.98% | 29.83m |
| DIMES | RL+AS+MCTS | 16.84 | 1.76% | 2.15h | 23.69 | 2.46% | 4.62h | 74.06 | 3.19% | 3.57h |
| OURS (DIFUSCO) | SL+MCTS | 16.63 | 0.46% | 10.13m | 23.39 | 1.17% | 24.47m | 73.62 | 2.58% | 47.36m |

06 Result

❖ Generalization Tests

- ✓ Results When Trained on TSP50 and Applied to Larger Instances



| METHOD | TYPE | SATLIB | | | ER-[700-800] | | |
|--------|------------|-----------------|------------------|-------------------|-----------------|------------------|-------------------|
| | | SIZE \uparrow | GAP \downarrow | TIME \downarrow | SIZE \uparrow | GAP \downarrow | TIME \downarrow |
| KAMIS | HEURISTICS | 425.96* | — | 37.58m | 44.87* | — | 52.13m |
| GUROBI | EXACT | 425.95 | 0.00% | 26.00m | 41.38 | 7.78% | 50.00m |
| INTEL | SL+G | 420.66 | 1.48% | 23.05m | 34.86 | 22.31% | 6.06m |
| INTEL | SL+TS | N/A | N/A | N/A | 38.80 | 13.43% | 20.00m |
| DGL | SL+TS | N/A | N/A | N/A | 37.26 | 16.96% | 22.71m |
| LWD | RL+S | 422.22 | 0.88% | 18.83m | 41.17 | 8.25% | 6.33m |
| DIMES | RL+G | 421.24 | 1.11% | 24.17m | 38.24 | 14.78% | 6.12m |
| DIMES | RL+S | 423.28 | 0.63% | 20.26m | 42.06 | 6.26% | 12.01m |
| OURS | SL+G | 424.50 | 0.34% | 8.76m | 38.83 | 12.40% | 8.80m |
| OURS | SL+S | 425.13 | 0.21% | 23.74m | 41.12 | 8.36% | 26.67m |

07 Conclusion

❖ Contribution

- ✓ One of the first successful applications of diffusion models to combinatorial optimization (CO) problems
- ✓ Demonstrates stronger scalability, expressiveness, and performance compared to traditional solvers
- ✓ Achieves efficient and generalizable performance on both TSP and MIS tasks, outperforming prior approaches

❖ Future Work

- ✓ Extension to broader NP-Complete (NPC) problems
- ✓ Integration of Equivariant Graph Neural Networks (GNNs)
- ✓ Exploration of accelerated inference techniques for faster sampling